

# Writeup for Q1

Steps followed for compiling the kernel and adding system call:

1. Download kernel source code using wget command

wget

<https://www.kernel.org/pub/linux/kernel/v4.x/linux-4.17.4.tar.xz>

2. Extract the tar.gz file using

`sudo tar -xvf linux-4.17.4.tar.xz -C/usr/src/`

3. Goto to directory kernel/ in the extracted folder

`cd /usr/src/linux-4.17.4/`

4. Edit file sys.c, ex : `sudo nano sys.c`

```
*/
SYSCALL_DEFINE2(mod_uruntime, int, p_pid, long, time )
{
    pid_t pid = p_pid;
    struct pid *pid_struct;
    struct task_struct *task;

    pid_struct = find_get_pid(pid);
    task = pid_task(pid_struct, PIDTYPE_PID);

    if(task == NULL){
        printk("Process not found\n");
        return -ESRCH;
    }

    task->se.uruntime += time;
    printk("Process %d's uruntime changed successfully ",p_pid);

    return 0;
}
```

5. Add your syscall code in the file sys.c

6. Go to directory arch/x86/entry/syscalls/ in the extracted folder

7. Add your system call to file syscall\_64.tbl

Ex: 440 common sample sys\_sample

```
547    x32    pwritev2                compat_sys_pwritev64v
548    common mod_vruntime            sys_mod_vruntime
# This is the end of the legacy x32 range. Numbers 548 and above are
# not special and are not to be used for x32-specific syscalls.
-- INSERT --
```

8. Edit config file so that only required things are compiled  
sudo make menuconfig

9. Compile the kernel with sudo make -j(\$number of processors to use)  
sudo make -jn

10 .Install modules using : sudo make modules\_install  
install  
sudo make modules\_install install

11. Reboot system  
sudo reboot

**Name of my syscall :- mod\_vruntime();**  
**short for modify runtime**

## The output of the test.c file

```
[kern@artixcse231 ~]$ gcc test.c -o test
[kern@artixcse231 ~]$ ./test
This is parent process :- 2679
It is running with mod_vruntime system call
And the time taken to complete its execution is:-0.000426
This is child process :- 2680
It is running without mod_vruntime system call
And the time taken to complete its execution is:-0.000286
[kern@artixcse231 ~]$
```

### Explanation: -

I have created two processes using fork() i.e child and parent. The child process is executing without mod\_vruntime syscall and the parent process is running with mod\_vruntime.

**We can clearly see that child process took:- 0.000286 sec  
While parent process took:- 0.000426 sec**

As we know that process having min vruntime are selected by the scheduler first in the CFS algorithm.

If we increase the vruntime of the process using sched\_entity structure then we can delay its selection because when we increase the vruntime then there will be self-balancing in the RB-tree and its execution time will be delayed.

**Data structures used are :- struct in c**

### **To achieve this:-**

The definition of struct sched\_entity is in file /include/linux/sched.h.

Now we change the actual functionality of the CFS scheduler. The CFS scheduler used a red-black tree to store and retrieve information about all the processes. By default the red-black tree picks the process with minimum vruntime.

This is achieved by changing the implementation of the function update\_curr() in file kernel/sched/fair.c.

Now the next step is defining the System Call called rt\_nice to change the value of vruntime of a process. In the system call all the information about the process is obtained using its pid in a task\_struct using pid\_task() and find\_get\_pid().

Error Handling Done in System Call :

2. Handling the case when the process is not found Code :

```
if (task == NULL) {  
    printk("Process not found\n");  
    return -ESRCH;  
}
```

## Relevant structures used:-

```
struct sched_entity {
    struct load_weight  load;          /* for load-balancing */
    struct rb_node      run_node;
    struct list_head    group_node;
    unsigned int        on_rq;

    u64                 exec_start;
    u64                 sum_exec_runtime;
    u64                 vruntime;
    u64                 prev_sum_exec_runtime;

    u64                 nr_migrations;

#ifdef CONFIG_SCHEDSTATS
    struct sched_statistics statistics;
#endif

#ifdef CONFIG_FAIR_GROUP_SCHED
    struct sched_entity *parent;
    /* rq on which this entity is (to be) queued: */
    struct cfs_rq        *cfs_rq;
    /* rq "owned" by this entity/group: */
    struct cfs_rq        *my_q;
#endif
};
```

Sched\_entity defination